

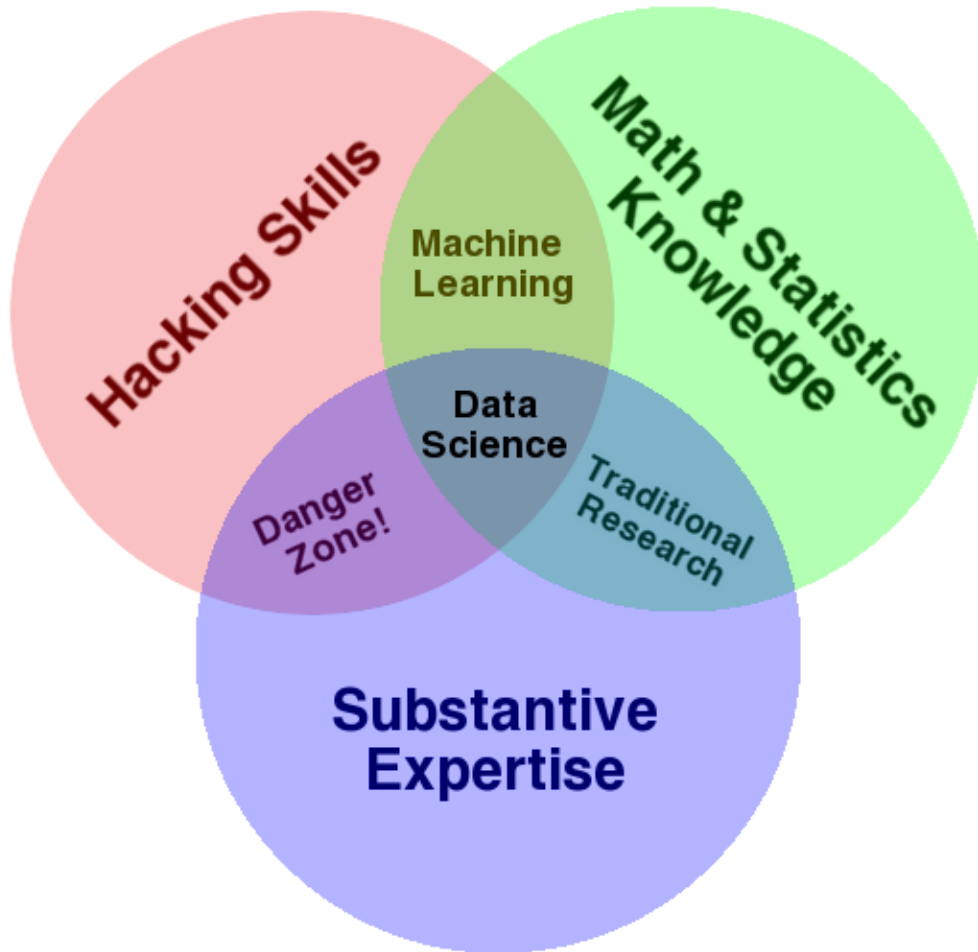
**Everything you ever
wanted to know about
data science (but were
too afraid to ask)**

Chris Beeley

What is data science?

- “A data scientist knows more about computer science than the average statistician, and more about statistics than the average computer scientist”

Drew Conway's famous Venn diagram



Source

Around the web...

- The difference between a statistician and a data scientist? About \$30,000
- ... an actual definition of data science. Taking a database and making it do something else. (warning: this quote is me! 😏)
- Statistics done on a Mac



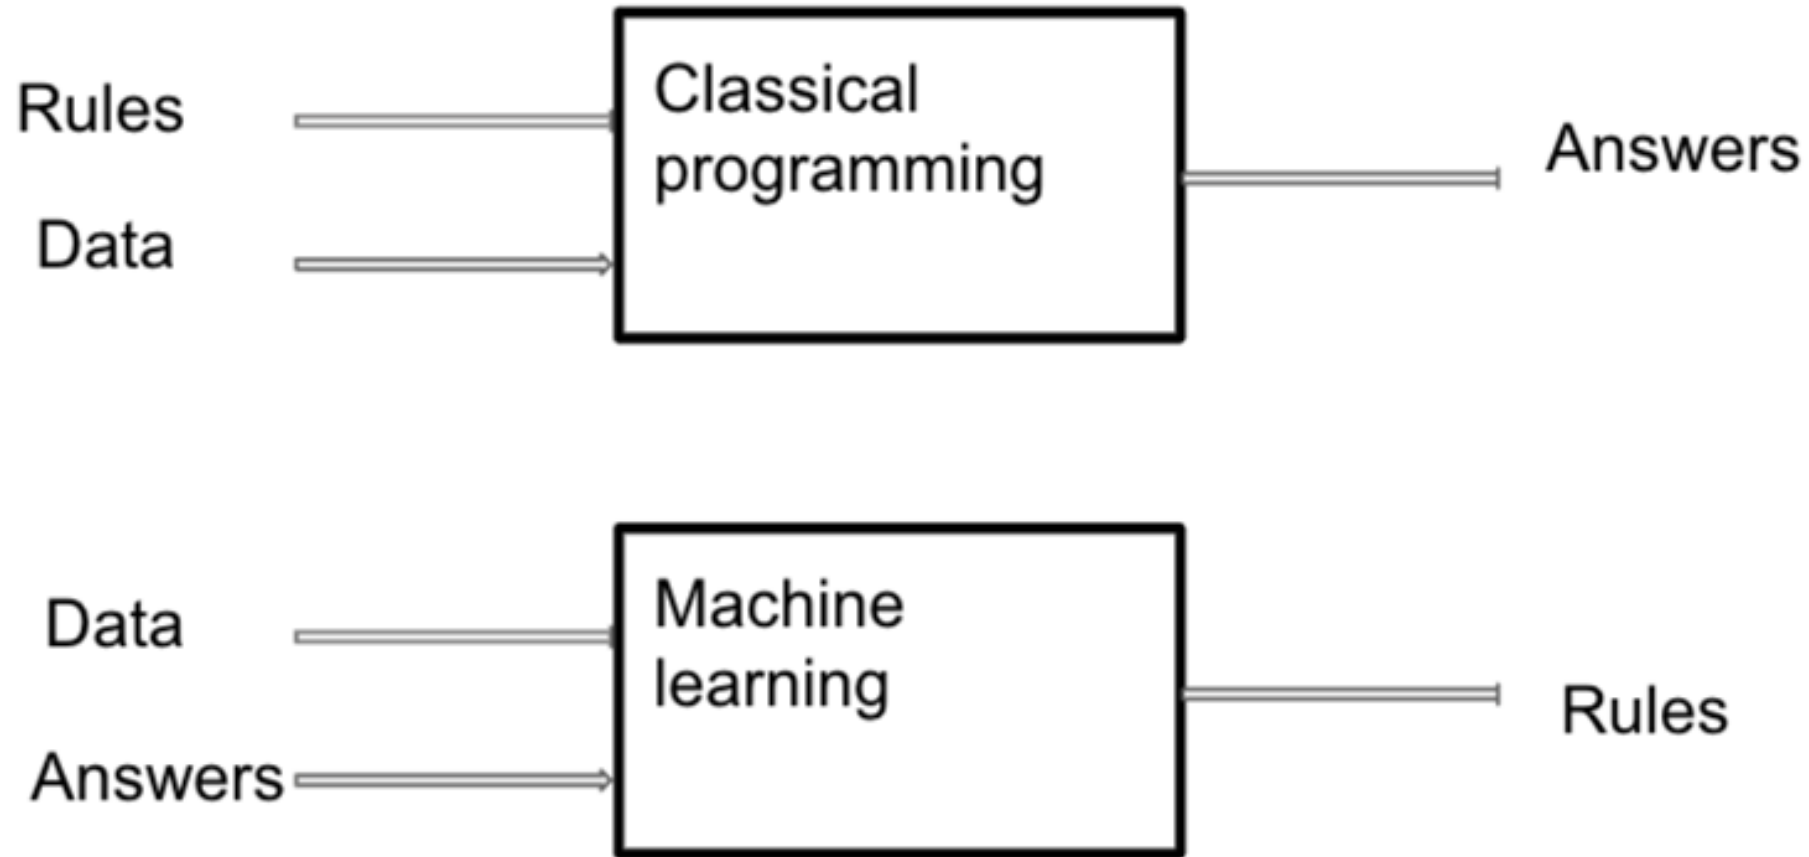
What are the skills of data science?

- Analysis
 - ML
 - Stats
 - Data viz
- Software engineering
 - Programming
 - SQL/ data
 - DevOps
 - RAP

What are the skills of data science?

- Domain knowledge
 - Communication
 - Problem formulation
 - Dashboards and reports

ML



Source

Inevitable XKCD



- Google flu trends

Source

Stats and data viz

- ML leans a bit more towards atheoretical prediction
- Stats leans a bit more towards inference (but they both do both)
- Data scientists may use different visualisations
 - Interactive web based tools
 - Dashboard based visualisers e.g. [stminsights](#)

Software engineering

- Programming
 - No/ low code data science?
- SQL/ data
 - Tend to use reproducible automated processes
- DevOps
 - Plan, code, build, test, release, deploy, operate, monitor
- RAP
 - I will come back to this

Domain knowledge

- Do stuff that matters
 - The best minds of my generation are thinking about how to make people click ads. That sucks. [Jeffrey Hammerbacher](#)
- Convince other people that it matters
- This is the hardest part of data science
- Communicate, communicate, communicate!
- Many of you are expert at this

Reproducibility

- Reproducibility in science
- The \$6B spreadsheet error
- George Osbourne's austerity was based on a spreadsheet error
- For us, reproducibility also means we can do the same analysis 50 times in one minute
 - Which is why I started down the road of data science

What is RAP

- a process in which code is used to minimise manual, undocumented steps, and a clear, properly documented process is produced in code which can reliably give the same result from the same dataset
- RAP should be:
 - the core working practice that must be supported by all platforms and teams; make this a core focus of NHS analyst training

Goldacre review

Levels of RAP- Baseline

- Data produced by code in an open-source language (e.g., Python, R, SQL).
- Code is version controlled (see Git basics and using Git collaboratively guides).
- Repository includes a README.md file (or equivalent) that clearly details steps a user must follow to reproduce the code
- Code has been peer reviewed.
- Code is published in the open and linked to & from accompanying publication (if relevant).

Levels of RAP- Silver

- Code is well-documented...
- Code is well-organised following standard directory format
- Reusable functions and/or classes are used where appropriate
- Pipeline includes a testing framework
- Repository includes dependency information (e.g. requirements.txt, PipFile, environment.yml)
- Data is handled and output in a [Tidy data](#) format

Source: [NHS Digital RAP community of practice](#)

Levels of RAP- Gold

- Code is fully packaged
- Repository automatically runs tests etc. via CI/CD or a different integration/deployment tool e.g. GitHub Actions
- Process runs based on event-based triggers (e.g., new data in database) or on a schedule
- Changes to the RAP are clearly signposted. E.g. a changelog in the package, releases etc. (See gov.uk info on Semantic Versioning)

Source: NHS Digital RAP community of practice

The data science “Unicorn”

- The maybe-mythical data science “Unicorn” has mastered:
 - Domain knowledge
 - Stats and ML
 - Software engineering

Data science is a team sport

- In my extended DS team I have:
- Stats and DevOps (and rabble rousing) [this one is me 🤔]
- SQL, data, and training
- DevOps and programming
- Text mining, Python, and APIs
- Bilingual R/ Python, Shiny dashboards

Data science is an MMO

- Data scientists need help with:
 - Stakeholder communication and engagement
 - Qualitative analysis
 - Translating models and prediction into the real world
 - Evidence review and problem definition

Data science is an MMO

- Data scientists are an excellent help when you:
 - Need a lot of pretty graphs
 - Need the same analysis done 50+ times with different data
 - Have too much text and not enough time to analyse it
 - Want to carefully document your analysis and make it reproducible
 - Have a hideously messy, large dataset that you can't hack together yourself

The team

Note

All copyrighted material is reused under [Fair Dealing](#)

